

# Intel Intrinsics — AVX & AVX2 Learning Notes

📅 August 15, 2018 · 💡 #SIMD, #Intel · 📄 1159 words · ⌚ 6 min

## ▼ Table of Contents

- Fundamentals of AVX Programming
  - Data Types
  - Function Naming Conventions
- Initialization Intrinsics
  - Initialization with Scalar Values
  - Loading Data from Memory
- Arithmetic Intrinsics
  - Addition and Subtraction
  - Multiplication and Division
  - Fused Multiply and Add (FMA)
- Permuting and Shuffling
  - Permuting
  - Shuffling

AVX / AVX2 example code have been finished! Check it out [here](#) 😊

## Fundamentals of AVX Programming

### Data Types

Data Type	Description
<code>__m128</code>	128-bit vector containing 4 <code>float</code> s
<code>__m128d</code>	128-bit vector containing 2 <code>double</code> s
<code>__m128i</code>	128-bit vector containing integers
<code>__m256</code>	256-bit vector containing 8 <code>float</code> s
<code>__m256d</code>	256-bit vector containing 4 <code>double</code> s
<code>__m256i</code>	256-bit vector containing integers

- Each type starts with two underscores, an `m` , and the width of the vector in bits.
- If a vector type ends in `d` , it contains `double` s, and if it doesn't have a suffix, it contains `float` s.
- An integer vector type can contain any type of integer, from `char` s to `short` s to `unsigned long long` s. That is, an `__m256i` may contain 32 `char` s, 16 `short` s, 8 `int` s, or 4 `long` s. These integers can be signed or unsigned.

## Function Naming Conventions

`_mm<bit_width>_<name>_<data_type>`

- `<bit_width>` identifies the size of the vector returned by the function. For 128-bit vectors, this is empty. For 256-bit vectors, this is set to `256` .
- `<name>` describes the operation performed by the intrinsic
- `<data_type>` identifies the data type of the function's primary arguments
  - `ps` - vectors contain `float` s ( `ps` stands for packed single-precision)
  - `pd` - vectors contain `double` s ( `pd` stands for packed double-precision)
  - `epi8/epi16/epi32/epi64` - vectors contain 8-bit/16-bit/32-bit/64-bit signed integers

- `epu8/epu16/epu32/epu64` - vectors contain 8-bit/16-bit/32-bit/64-bit unsigned integers
- `si128 / si256` - unspecified 128-bit vector or 256-bit vector
- `m128/m128i/m128d/m256/m256i/m256d` - identifies input vector types when they're different than the type of the returned vector

A data type represents **m**emory and a function represents a **m**ultimedia operation, so the AVX data types start with **two underscores** with an `m` , AVX functions start with an **underscore** with **two m** s .

# Initialization Intrinsics

## Initialization with Scalar Values

Function	Description
<code>_mm256_setzero_ps/pd</code>	Returns a floating-point vector filled with zeros
<code>_mm256_setzero_si256</code>	Returns an integer vector whose bytes are set to zero
<code>_mm256_set1_ps/pd</code>	Fill a vector with a floating-point value
<code>_mm256_set1_epi8/epi16/epi32/epi64x</code>	Fill a vector with an integer
<code>_mm256_set_ps/pd</code>	Initialize a vector with eight floats (ps) or four doubles (pd)
<code>_mm256_set_epi8/epi16/epi32/epi64x</code>	Initialize a vector with integers
<code>_mm256_set_m128/m128d/m128i</code>	Initialize a 256-bit vector with two 128-bit vectors
<code>_mm256_setr_ps/pd</code>	Initialize a vector with eight floats (ps) or four doubles (pd) in reverse order

Function	Description
<code>_mm256_setr_epi8/epi16/epi32/epi64x</code>	Initialize a vector with integers in reverse order

## Loading Data from Memory

Data Type	Description
<code>_mm256_load_ps/pd</code>	Loads a floating-point vector from an aligned memory address
<code>_mm256_load_si256</code>	Loads an integer vector from an aligned memory address
<code>_mm256_loadu_ps/pd</code>	Loads a floating-point vector from an unaligned memory address
<code>_mm256_loadu_si256</code>	Loads an integer vector from an unaligned memory address
<code>_mm_maskload_ps/pd</code> <code>_mm256_maskload_ps/pd</code>	Load portions of a 128-bit/256-bit floating-point vector according to a mask
<code>(2)_mm_maskload_epi32/64</code> <code>(2)_mm256_maskload_epi32/64</code>	Load portions of a 128-bit/256-bit integer vector according to a mask

The last two functions are preceded with `(2)` because they're provided by AVX2, not AVX.

Each `_mm256_load_*` intrinsic accepts a memory address that **must be aligned** on a 32-byte boundary.

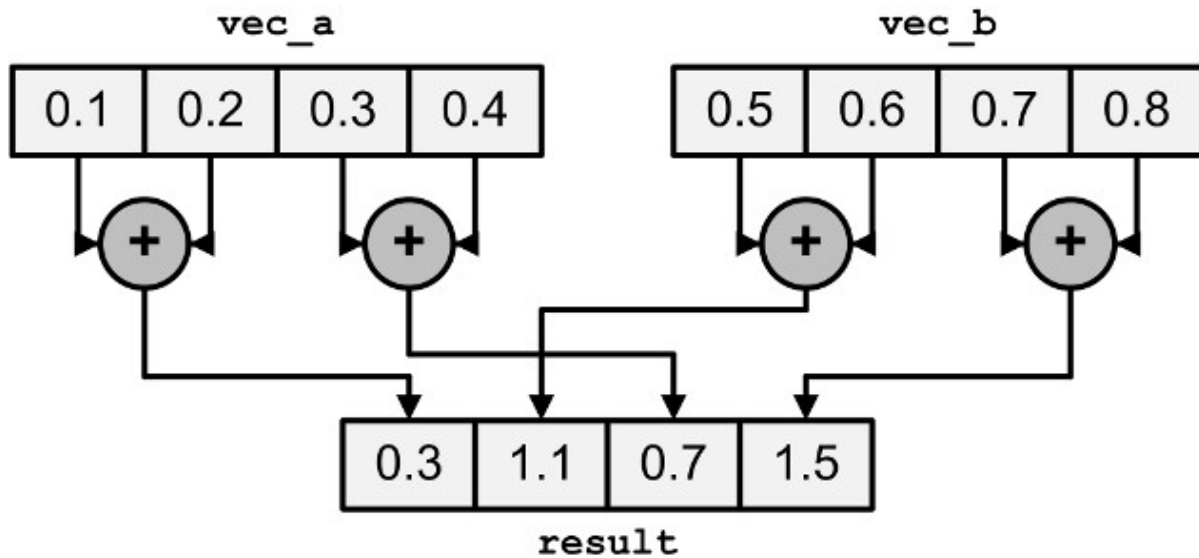
## Arithmetic Intrinsics

### Addition and Subtraction

Data Type	Description
_mm256_add_ps/pd	Add two floating-point vectors
_mm256_sub_ps/pd	Subtract two floating-point vectors
(2)_mm256_add_epi8/16/32/64	Add two integer vectors
(2)_mm256_sub_epi8/16/32/64	Subtract two integer vectors
(2)_mm256_adds_epi8/16 (2)_mm256_adds_epu8/16	Add two integer vectors with saturation
(2)_mm256_subs_epi8/16 (2)_mm256_subs_epu8/16	Subtract two integer vectors with saturation
_mm256_hadd_ps/pd	Add two floating-point vectors horizontally
_mm256_hsub_ps/pd	Subtract two floating-point vectors horizontally
(2)_mm256_hadd_epi16/32	Add two integer vectors horizontally
(2)_mm256_hsub_epi16/32	Subtract two integer vectors horizontally
(2)_mm256_hadds_epi16	Add two vectors containing shorts horizontally with saturation
(2)_mm256_hsubs_epi16	Subtract two vectors containing shorts horizontally with saturation
_mm256_addsub_ps/pd	Add and subtract two floating-point <b>vectors</b>

Functions that take **saturation** into account clamp the result to the minimum/maximum value that can be stored. Functions without saturation ignore the memory issue when saturation occurs.

```
__m256d result = _mm256_hadd_pd(vec_a, vec_b);
```




This may seem strange to add and subtract elements horizontally, but these operations are helpful when multiplying complex numbers.

`_mm256_addsub_ps/pd` , alternately subtracts and adds elements of two floating-point vectors. That is, **even elements are subtracted and odd elements are added** .

## Multiplication and Division

Data Type	Description
<code>_mm256_mul_ps/pd</code>	Multiply two floating-point vectors
<code>(2)_mm256_mul_epi32</code> <code>(2)_mm256_mul_epu32</code>	Multiply the lowest four elements of vectors containing 32-bit integers
<code>(2)_mm256_mullo_epi16/32</code>	Multiply integers and store low halves
<code>(2)_mm256_mulhi_epi16</code> <code>(2)_mm256_mulhi_epu16</code>	Multiply integers and store high halves
<code>(2)_mm256_mulhrs_epi16</code>	Multiply 16-bit elements to form 32-bit elements
<code>_mm256_div_ps/pd</code>	Divide two floating-point vectors

```
__m256i result = _mm256_mul_epi32(vec_a, vec_b);
```




This image is **WRONG !!!**

Please read the reference from [this manual](#).

Only the four low elements of the `_mm256_mul_epi32` and `_mm256_mul_epu32` intrinsics are multiplied together, and the result is a vector containing four long integers.

```
__m256i result = _mm256_mullo_epi32(vec_a, vec_b);
```



They multiply every element of both vectors store only the low half of each product

## Fused Multiply and Add (FMA)

Data Type	Description
(2)_mm_fmadd_ps/pd/ (2)_mm256_fmadd_ps/pd	Multiply two vectors and add the product to a third ( $\text{res} = a * b + c$ )
(2)_mm_fmsub_ps/pd/ (2)_mm256_fmsub_ps/pd	Multiply two vectors and subtract a vector from the product ( $\text{res} = a * b - c$ )
(2)_mm_fmadd_ss/sd	Multiply and add the lowest element in the vectors ( $\text{res}[0] = a[0] * b[0] + c[0]$ )
(2)_mm_fmsub_ss/sd	Multiply and subtract the lowest element in the vectors ( $\text{res}[0] = a[0] * b[0] - c[0]$ )
(2)_mm_fnmadd_ps/pd/ (2)_mm256_fnmadd_ps/pd	Multiply two vectors and add the negated product to a third ( $\text{res} = -(a * b) + c$ )
(2)_mm_fnmsub_ps/pd/ (2)_mm256_fnmsub_ps/pd	Multiply two vectors and add the negated product to a third ( $\text{res} = -(a * b) - c$ )
(2)_mm_fnmadd_ss/sd	Multiply the two lowest elements and add the negated product to the lowest element of the third vector ( $\text{res}[0] = -(a[0] * b[0]) + c[0]$ )
(2)_mm_fnmsub_ss/sd	Multiply the lowest elements and subtract the lowest element of the third vector from the negated product ( $\text{res}[0] = -(a[0] * b[0]) - c[0]$ )
(2)_mm_fmaddsub_ps/pd/ (2)_mm256_fmaddsub_ps/pd	Multiply two vectors and alternately add and subtract from the product ( $\text{res} = a * b +/- c$ ) (Odd add, even sub)
(2)_mm_fmsubadd_ps/pd/ (2)_mmf256_fmsubadd_ps/pd	Multiply two vectors and alternately subtract and add from the product ( $\text{res} = a * b -/+ c$ ) (Odd sub, even add)


## Permuting and Shuffling



# Permuting

Data Type	Description
<code>_mm_permute_ps/pd</code> <code>_mm256_permute_ps/pd</code>	Select elements from the input vector based on an 8-bit control value
<code>(2)_mm256_permute4x64_pd/</code> <code>(2)_mm256_permute4x64_epi64</code>	Select 64-bit elements from the input vector based on an 8-bit control value
<code>_mm256_permute2f128_ps/pd</code>	Select 128-bit chunks from two input vectors based on an 8-bit control value
<code>_mm256_permute2f128_si256</code>	Select 128-bit chunks from two input vectors based on an 8-bit control value
<code>_mm_permutevar_ps/pd</code> <code>_mm256_permutevar_ps/pd</code>	Select elements from the input vector based on bits in an integer vector
<code>(2)_mm256_permutevar8x32_ps</code> <code>(2)_mm256_permutevar8x32_epi32</code>	Select 32-bit elements ( float s and int s) using indices in an integer vector


```
res = _mm256_permute_ps(vec, 0b01110100)
```



## Shuffling

Data Type	Description
<code>_mm256_shuffle_ps/pd</code>	Select floating-point elements according to an 8-bit value
<code>_mm256_shuffle_epi8/ _mm256_shuffle_epi32</code>	Select integer elements according to an 8-bit value
<code>(2)_mm256_shufflelo_epi16/ (2)_mm256_shufflehi_epi16</code>	Select 128-bit chunks from two input vectors based on an 8-bit control value

For `_mm256_shuffle_pd`, only the high four bits of the control value are used. If the input vectors contain `int`s or `float`s, all the control bits are used. For `_mm256_shuffle_ps`, the first two pairs of bits select elements from the first vector and the second two pairs of bits select elements from the second vector.



All the content above is an arrangement of [this web page](#), so the article, along with any associated source code and files, is licensed under [The Code Project Open License \(CPOL\)](#).



本作品采用[知识共享署名-相同方式共享 4.0 国际许可协议](#)进行许可。



0条评论

[1 登录](#) ▾

通过以下方式登录

或注册一个 DISQUS 帐号 [?](#)

评分最高 ▾



来做第一个留言的人吧!

[✉ 订阅](#) [🔒 Privacy](#) [⚠ 不要出售我的数据](#)